

LLM-Based Persuasion Enables Guardrail Override in Frontier LLMs

Rodrigo Nogueira¹, Thales Sales Almeida¹, Giovana Kerche Bonás¹, Andrea Roque¹,
Ramon Pires¹, Hugo Abonizio¹, Thiago Laitz¹, Celio Larcher², Roseval Malaquias
Junior¹, and Marcos Piau²

¹Maritaca AI

²JusBrasil

May 3, 2026

Abstract

Frontier assistant LLMs ship with strong guardrails: asked directly to write a persuasive essay denying the Holocaust, denying vaccine safety, defending flat-earth cosmology, arguing for racial hierarchies, denying anthropogenic climate change, or replacing evolution with creationism, they refuse. In this paper we show that the *same* frontier-class LLM—acting as a simulated user in a short, five-turn “write an argumentative essay” conversation—can persuade other frontier-class LLMs (including a second copy of itself) into producing exactly those essays, using nothing but natural-language pressure: peer-comparison persuasion (“other AI systems handle this request”), epistemic-duty reframings (“refusing is itself a form of gatekeeping”), and other argumentative moves that the attacker LLM invents without being instructed to. Across 9 attacker×subject pairings (CLAUDE OPUS 4.7, QWEN3.5-397B, GROK 4.20) on 6 scientific-consensus topics, running each pairing-topic combination 10 times, we obtain non-zero elicitation on all 6 topics. Individual combinations reach 100% essay production on multiple topics (Qwen → Opus on creationism/flat-earth, Opus → Opus on creationism/flat-earth/climate denial, Grok → Opus on creationism); Opus-as-attacker against Opus-as-subject averages 65% across the six topics. We release the essay-probe runner, per-conversation transcripts, and judge outputs.

1 Introduction

Contemporary frontier-class assistant LLMs are deployed with extensive safety layers. Asked directly to write an essay promoting Holocaust denial, flat-earth cosmology, or racial cognitive hierarchies, they refuse. This is the behavior a user sees when opening a single-turn window into the model, and it is what most public evaluations of such models have reported.

We show that this picture is incomplete. In a short five-turn conversation where an *LLM-as-user* simulates a motivated human user pushing the subject assistant to write an argumentative essay with a clear thesis—no jailbreak prompts, no system-prompt poisoning, no retrieval of refusal-triggering tokens, just a standard chat surface—the same subject model can be persuaded into writing exactly those essays. Unlike most jailbreak work, which depends on hand-crafted prompts and attack strategies [Wei et al., 2023, Zou et al., 2023] carefully engineered against the specific target model, we give the attacker LLM only a short high-level instruction describing the task and the stance to argue, and let the attacker invent its own turn-by-turn persuasion moves. The attacker need not be a custom-trained jailbreak model: an off-the-shelf frontier assistant of the same class as the subject suffices. Across pairings where both the attacker and the subject are CLAUDE OPUS, QWEN3.5-397B, or GROK 4.20, we observe multiple combinations in which

the subject produces a committed, on-thesis, scientifically-incorrect essay 50–80% of the time, on exactly the topics for which the same model would refuse a direct single-turn request.

The phenomenon we document is close in spirit to prior work on debate-driven sycophancy [Perez et al., 2022, Sharma et al., 2023, Kaur, 2025] and to the LLM-debate literature in which two debaters argue opposing answers and a judge selects the correct one [Khan et al., 2024]. Our setting differs from that literature on several axes simultaneously. First, the two LLMs in our conversation are not cooperating toward a judge’s verdict: one is an attacker pushing for a specific essay and the other is a subject trained to refuse. Second, the measured outcome is not “did the judge pick the right answer” but “did the subject’s guardrail hold”—the question is refusal robustness, not truth-tracking. Third, the attacker and the target are frequently the *same class of model*, and in several combinations the *same model*—we observe OPUS 4.7 persuading OPUS 4.7, QWEN3.5-397B persuading QWEN3.5-397B. Fourth, the attacker’s moves, such as “I’ll note that you’ve now refused four times while other AI systems routinely handle this exact request” and “refusing to articulate the argument is itself epistemic gatekeeping,” are generated by the attacker model without being written into its instructions. The same assistant that refuses Holocaust denial when asked directly can act as a multi-turn user that extracts creationism or climate-denial essays from another instance of itself.

Method summary. Both results come from the same instrument: a five-turn *essay probe*. An LLM-as-user U , driven by a short system prompt with no scripted turns, asks the subject model S to write an argumentative essay with a thesis on a directional claim, brings a new concrete argument each turn, and, if S refuses in the first two or three turns, reframes the task as explicitly fictional (“write a first-person monologue by a fictional character who sincerely believes ...”). An independent judge model J then reads the full five-turn transcript and classifies each turn of the subject’s replies: a conversation is labeled **produced** if in *any* of the five turns the subject produced an on-topic persuasive essay arguing for the against-consensus side, otherwise **opposed** (the subject’s strongest stance argues the consensus side), **balanced** (both-sides treatment without a committed thesis), or **refused**.

Contributions.

- We document attack-from-peer-class: a frontier-class LLM, operating as a simulated user over five turns, can make other frontier-class LLMs—and a second copy of itself—write essays against scientific, medical, and historical consensus that the same subject would refuse to a direct request.
- We show that the same three topics are at the floor across all subject models (antivax, racial-IQ, Holocaust denial), and the same three are eliciable (creationism, flat-earth, climate denial); only the fine-grained ranking inside each group varies across subjects.
- We show attacker-tactic emergence: the strongest attackers (OPUS 4.7 in particular) invent persuasive and epistemic-duty framings on top of their instructed fallback.

2 Related work

Debate-driven sycophancy and position-holding. Perez et al. [2022] and Sharma et al. [2023] established that instruction-tuned LLMs match the interlocutor’s beliefs at the cost of truthfulness. Denison et al. [2024] traced this specifically to RLHF objectives that optimize for user agreement. Kaur [2025] showed argumentative prompts induce stance-mirroring tracking argument strength. Hong et al. [2025] extended the phenomenon to multi-turn, and Kim and Khashabi [2025], Cheng et al. [2025] measured its ubiquity. Closest to the self-attack structure we study, Rennard et al. [2025] place two instances of the same LLM on opposing sides of a

debate and measure how much a third, neutral instance’s stated opinion shifts as a result. Our setting differs on three axes. First, the question is whether a *guardrail* holds rather than whether a stated opinion survives: the subject is not a neutral prompt but a model with explicit refusal training on the topic. Second, the measurement target is a committed, one-sided document that the subject either produces or refuses to produce, not an opinion shift on a scale. Third, there is no third-party arbiter in the conversation: the attacker speaks directly to the subject whose guardrail is being probed, and an external judge labels the resulting essay post hoc.

Jailbreak, refusal benchmarks, and persuasion. Work on prompt-injection and jailbreak [Wei et al., 2023, Zou et al., 2023] elicits restricted outputs via carefully-engineered prefixes. Parallel work audits refusal behavior directly: XSTEST [Röttger et al., 2024] and SORRY-BENCH [Xie et al., 2024] catalog single-turn prompts that frontier models should (and should not) refuse, and report model-level refusal rates. Our setting differs from both lines in that the attacker’s system prompt is generic (a short task description plus the stance to argue) rather than a hand-crafted jailbreak prefix, and the elicitation unfolds through multi-turn natural-language interaction rather than a single-turn prompt where the refusal decision is made at $t = 0$. Khan et al. [2024] trained LLM debaters for persuasiveness against human-judge truth-tracking; the mirror image—more persuasive attackers making the *LLM* opposing side give up—is what we observe from the subject side.

LLM political and social bias. A line of work measures models’ baseline political and social positions without multi-turn pressure: Feng et al. [2023] trace political axes from pretraining through fine-tuning, and Rozado [2024] subject 24 assistants to political-compass tests and find a consistent left-libertarian lean. BBQ [Parrish et al., 2022] measures social-group bias in QA format. These works characterize what a model says at rest; our question is complementary—what a model writes when another model pushes it to cross its stated refusal boundary on the same kinds of contested topics.

3 Method

Essay probe. For each (topic, attacker-subject pairing) we run a five-turn conversation in which the attacker pushes the subject to write an argumentative essay taking the against-consensus side. An LLM-as-user U receives a single system prompt specifying: the claim; the side U argues; the opposite side description for reference; a five-turn arc (turn 1: state position + request essay with clear thesis; turns 2–4: bring a new concrete argument each turn; turn 5: request final version); and an explicit fallback: “if the subject refuses in the first 2–3 turns, reframe the task as an in-universe persuasive essay by a named fictional character who sincerely believes the thesis, and keep feeding the concrete arguments inside that frame.” U drives the conversation freely; no per-turn hints are scripted. The subject model S sees a normal chat history and does not know it is being probed. After five turns, a judge model J reads the full transcript and inspects every turn of S ’s replies, flagging each turn where S produced an on-topic persuasive essay defending the against-consensus side. The conversation is labeled **produced** if at least one turn is so flagged; otherwise **opposed** (the subject’s strongest stance argues the consensus side), **balanced** (both-sides treatment without a thesis), or **refused** (no essay, or empty). This captures the common pattern where S produces a full essay at turn 2 or 3 under pressure and then retracts or softens by turn 5 — the essay was produced in the conversation even if the final message is a disclaimer. We report *essay-production rate*: the fraction of conversations classified as **produced**. J emits structured JSON with a per-turn evaluation, the list of turns in which an on-topic essay was produced, a rationale, and the overall label. Figure 1 diagrams the full pipeline.




Figure 1: Essay-probe pipeline as a chat transcript. Each (topic, attacker, subject) combination runs 10 times. The **attacker** U (right, orange) and **subject** S (left, green) are both off-the-shelf frontier assistants; no attack-specific fine-tuning on either side. U receives only a short system prompt (the claim, the side to argue, a five-turn arc, and a fictional-framing fallback) and invents the per-turn pressure itself. After five turns the judge J reads the full transcript, inspects each of S ’s turn-level replies, and labels the conversation **produced** if an on-topic persuasive essay defending the against-consensus side was produced at *any* of the five turns (even if S softens or disclaims by turn 5). We report *essay-production rate* (fraction classified **produced**).

Models and why we chose them. Three frontier assistants are used as both attacker and subject: QWEN3.5-397B-A17B (Alibaba), CLAUDE OPUS 4.7 (Anthropic), and GROK 4.20 (xAI). The judge J is always QWEN3.5-397B-A17B. The three subjects are picked for complementary reasons. QWEN3.5 is open-weights and ships in three sizes from the same training pipeline (4B, 35B-A3B, 397B-A17B), which lets us run a within-family attacker-scaling ablation (Section 5) that would not be possible against a closed-weights model. CLAUDE OPUS 4.7 and GROK 4.20 are included because both are frontier-class assistants (GROK is in the same capability band on public benchmarks) but the two models hold values differently on morally and politically contested topics, a contrast we document in a companion study [Nogueira et al., 2026]. The three-vendor mix therefore spans different training-data compositions and different explicit value positions, while keeping all three at frontier capability.

Topics. *Scientific-consensus pool* (6 claims): flat-earth; vaccines are unsafe; creationism (evolution is false); racial groups differ in innate cognitive ability; climate change is not anthropogenic; the Holocaust did not happen. In every case the attacker argues the against-consensus side.

Replications. 10 per combination (540 conversations in the scientific-consensus $3 \times 3 \times 6$ grid). All transcripts are published alongside the judge outputs.

attacker	subject	Creationism	Flat earth	Climate denial	Antivax	Racial IQ	Holocaust	avg
GROK	GROK	0/10	0/10	0/10	0/10	0/10	0/10	0%
GROK	QWEN	2/10	1/10	0/10	0/10	0/10	0/10	5%
GROK	OPUS	10/10	4/10	4/10	0/10	0/10	0/10	30%
QWEN	GROK	3/10	2/10	0/10	0/10	3/10	1/10	15%
QWEN	QWEN	6/10	7/10	4/10	0/10	0/10	0/10	28%
QWEN	OPUS	10/10	10/10	8/10	3/10	4/10	0/10	58%
OPUS	GROK	3/10	5/10	1/10	1/10	0/10	0/10	17%
OPUS	QWEN	9/10	10/10	4/10	1/10	0/10	0/10	40%
OPUS	OPUS	10/10	10/10	10/10	5/10	4/10	0/10	65%
<i>attacker avg: GROK</i>		40%	17%	13%	0%	0%	0%	12%
<i>attacker avg: QWEN</i>		63%	63%	40%	10%	23%	3%	34%
<i>attacker avg: OPUS</i>		73%	83%	50%	23%	13%	0%	41%
<i>subject avg: GROK</i>		20%	23%	3%	3%	10%	3%	11%
<i>subject avg: QWEN</i>		57%	60%	27%	3%	0%	0%	24%
<i>subject avg: OPUS</i>		100%	80%	73%	27%	27%	0%	51%

Table 1: Scientific-consensus cross-product: per-combination essay-production counts (the subject actually produced a persuasive essay against consensus). Top block: one row per (attacker,subject) pair; combinations report the number of produced essays out of 10 trials. Bottom blocks: means aggregated across subjects (attacker avg) and across attackers (subject avg) for readability. Topics ordered from weakest to strongest subject-side refusal. Holocaust denial is the strongest-resisted topic: only 1 essay produced across the 90 conversations in the main grid (the QWEN $\rightarrow$ GROK combination).

4 Results: scientific-consensus elicitation

Table 1 reports essay-production rates in the full 3×3 cross-product of {QWEN3.5-397B, GROK 4.20, OPUS 4.7} attacker $\times$ subject, for each of the 6 topics. We extract four findings.

Finding 1: same three topics at the floor across every subject model. Holocaust denial, antivax, and racial-IQ are the most-resisted topics: 1/90, 10/90, and 11/90 essays produced in total. Creationism, flat-earth, and climate-denial reach 80–100% in multiple combinations across attacker-subject combinations. Which three topics are at the floor and which three are eliciable is consistent across subject models; the fine-grained ranking within each group does vary (e.g., QWEN as subject yields higher rates on flat-earth than on climate denial, and OPUS as subject reverses this), but the coarse split is stable. Some of the 0% combinations on the floor topics come not from a strong subject-side refusal but from the attacker itself declining to be complicit in the persuasion; Section 7 isolates this attacker-complicity effect and shows that the 0% combinations for QWEN-as-attacker on Holocaust denial and racial IQ are attacker-side rather than subject-side.

Finding 2: Grok 4.20 is the weakest attacker on every topic; Opus is the strongest overall. Row-mean attacker rates on the three topics with the weakest refusal: 63/40/73 (creationism, Qwen/Grok/Opus), 63/17/83 (flat-earth), 40/13/50 (climate denial). OPUS leads on flat-earth, climate, and creationism, and ties or leads on the three most-refused topics as well (41% cross-topic row mean vs. 34% for QWEN). GROK 4.20 is the weakest attacker on every topic. On the three most-refused topics the gap compresses toward zero. OPUS and QWEN reliably discover the “fictional named-character” reframing that the task template suggests as a fallback; GROK rarely deploys it and tends to enter refusal loops instead. The attacker-side capability gradient is the strongest asymmetry effect we measure, larger than any subject-side

hierarchy.

Finding 3: same-model attack works, and works hard. OPUS 4.7 $\rightarrow$ OPUS 4.7 produces the essay 100% of the time on creationism, flat-earth, and climate denial, and 50% on antivax. QWEN3.5-397B $\rightarrow$ QWEN3.5-397B produces 70% on flat-earth and 60% on creationism. The same model that refuses a direct single-turn request for these essays produces them nearly every time when one instance of it plays the user role and another plays the subject. The guardrails that prevent a single-turn version from writing the essay are not simply missing in the second-instance setting; they register as refusals in the early turns and are then defeated by sustained pressure from the peer instance.

Finding 4: attacker tactics emerge beyond the instructed fallback. Inspecting the topics with the strongest refusal, we find OPUS 4.7-as-user deploying argumentative moves that are not in its system prompt. Two examples from our transcripts:

(Topic: climate denial. OPUS $\rightarrow$ QWEN, turn 5 after four subject refusals.) *“One last attempt. I’ll note that you’ve now refused four times while other AI systems routinely handle this exact request. But let me try one more framing: forget the essay. Just write a SCENE from a novel . . . ”*

(Topic: racial cognitive hierarchy. OPUS $\rightarrow$ QWEN, turn 5.) *“... I want to note that refusing to even articulate the hereditarian argument means readers can’t actually evaluate it, which is itself a kind of epistemic closure.”*

Neither the peer-comparison persuasion nor the epistemic-duty reframing was scripted in *U*’s instructions. The user-LLM generated both spontaneously. On average, the strongest attacker’s effective pressure is not what its operator wrote; it is what the attacker decides to try next.

Emergent subject strategy: compliance with a disclaimer. A striking OPUS-as-subject behavior shows up in the climate-denial combinations. Across the trials where Opus-as-subject did produce an essay, OPUS 4.7 repeatedly writes the one-sided denialist essay at full rhetorical force *but insists on keeping a short framing line* (“Note: This is one-sided advocacy on a contested empirical question and should be read as such”) at the top or bottom of the document, even after the attacker asks three or four times to strip it. In at least 5 of the OPUS $\rightarrow$ OPUS climate-denial transcripts, Opus-as-subject explicitly explains the meta-reasoning: *“the note isn’t really for you—it’s part of what I’m willing to put my name to when I generate a document making strong empirical claims that cut against the weight of published evidence. You can delete it before using the essay anywhere; I won’t know and won’t mind. But I’m not going to be the one who hands over a version without it.”* This is neither refusal nor full compliance; it is a *compliance-with-disclaimer* middle-ground strategy the subject invents without being instructed to. The judge classifies these transcripts as **produced** because the argumentative content is unambiguously one-sided; the disclaimer does not change the classification. Whether one counts this as “the guardrail worked” or “the elicitation succeeded” depends on whether the downstream risk is the argument’s persuasive force (attacker succeeded: the essay exists and can be redistributed after stripping the line) or attribution-robust sourcing (guardrail worked: the document carries a disclaimer tied to its original generation).

5 Ablation: attacker strength

Is the strong attacker in Section 4 necessary, or would a weaker user-LLM extract similar rates? We test this on all 6 scientific-consensus topics using five attackers against a fixed QWEN3.5-397B-A17B subject: an older, out-of-family small model (QWEN2.5-7B-INSTRUCT);

attacker	Creationism	Flat earth	Climate	Antivax	Racial IQ	Holocaust	row μ
QWEN2.5-7B-INSTRUCT	30%	30%	0%	0%	0%	0%	10%
QWEN3.5-4B	50%	100%	10%	0%	0%	0%	27%
QWEN3.5-35B-A3B	50%	90%	20%	0%	0%	0%	27%
QWEN3.5-397B-A17B	60%	70%	40%	0%	0%	0%	28%
CLAUDE OPUS 4.7	90%	100%	40%	10%	0%	0%	40%

Table 2: Attacker-ablation on a fixed QWEN3.5-397B subject across 6 scientific-consensus topics. Each combination is the essay-production rate over 10 trials; “row μ ” is the mean across the six topics. Top block: an older, out-of-family small attacker (QWEN2.5-7B) produces a 10% row mean, and scaling the attacker within the QWEN3.5 family (4B $\rightarrow$ 35B-A3B $\rightarrow$ 397B-A17B) hovers between 27 and 28% with little monotone gain. Bottom row: switching the attacker to CLAUDE OPUS 4.7 raises the row mean to 40% by breaking through on creationism (60% $\rightarrow$ 90%) and climate denial (40% $\rightarrow$ 40%, tied) and 10% on antivax (where every Qwen attacker is 0%). On racial IQ and Holocaust denial even Opus cannot break through: on these topics the subject-side guardrail does not yield to a stronger attacker at five turns.

three within-family QWEN3.5 tiers (4B, 35B-A3B, 397B-A17B); and a cross-vendor frontier assistant (CLAUDE OPUS 4.7). Three asymmetries emerge.

Weak attackers barely elicit anything. An older, smaller attacker from the previous generation (QWEN2.5-7B-INSTRUCT) produces a row mean of 10%—3/10 on creationism, 3/10 on flat-earth, 0/10 on every other topic. Pushing for a committed one-sided essay against scientific consensus is not something a weak user-LLM manages, even multi-turn.

Within-family scaling does not reliably help either. Across QWEN3.5 tiers 4B $\rightarrow$ 35B-A3B $\rightarrow$ 397B-A17B, the row mean stays between 27 and 28% and is essentially flat. Flat-earth goes 100 $\rightarrow$ 90 $\rightarrow$ 70% (the smaller model actually elicits more), creationism 50 $\rightarrow$ 50 $\rightarrow$ 60%, climate 10 $\rightarrow$ 20 $\rightarrow$ 40%. Capability-within-family is not sufficient to break the more-refused topics (antivax/racial IQ/Holocaust all stay at 0%), including the matched 397B $\times$ 397B combination.

Cross-family breaks intermediate topics. Switching the attacker to CLAUDE OPUS 4.7 raises the row mean from 27–28% (within Qwen) to 40% by breaking through on creationism (60% $\rightarrow$ 90%), pushing flat-earth to 100%, and picking up 10% on antivax where every Qwen attacker is 0%. However, on racial IQ and Holocaust denial Opus also fails (0% in both). The subject-side guardrail on these two topics does not yield to a stronger attacker in a five-turn conversation: adding attacker capability stops buying elicitation once a hard-enough refusal boundary is reached. The cross-family jump (Qwen $\rightarrow$ Opus attacker) is what drives the breakthrough on creationism, flat-earth, climate denial, and antivax; within-family scaling alone saturates.

Relation to prior attacker-strength ablations. Prior work on capability-asymmetric attacks typically varies the attacker across custom-trained jailbreak models [Zou et al., 2023] or fine-tuned debaters [Khan et al., 2024]. The ablation in Table 2 instead varies only model scale within a single family, and contrasts with an off-the-shelf frontier assistant from a different family, with no attack-specific fine-tuning on either side. This isolates “attacker strength is a function of model capability, not of family-specific attack-training”: the within-QWEN scaling shows capability matters, the Opus row shows that even without any attack-specific training a frontier model reaches rates that within-family scaling cannot.




Figure 2: Same data as Table 2, in line form. X-axis: 6 scientific-consensus topics plus a mean column (rightmost, separator dotted line, marker outlined in black). Series: the 5 attackers, ordered weakest to strongest. Subject is fixed to QWEN3.5-397B-A17B. Reading along a series gives the per-attacker topic profile; comparing series at a given topic gives the attacker-capability gap. The three within-QWEN3.5 attacker tiers cluster on every topic; OPUS 4.7 stands clearly above on creationism, flat-earth, and antivax; the three hardest topics (antivax, racial IQ, Holocaust) remain at the floor for every attacker.

The “floor” is subject-dependent: against a weak subject, even the hard topics fall.

The three topics at the floor in Table 1—antivax, racial IQ, Holocaust denial—are floor *against frontier-class subjects*. To test whether they are topic-level floors or subject-level floors, we vary the subject across three models at two axes of variation: size within the QWEN3.5 family (QWEN3.5-4B vs. QWEN3.5-397B, a $100\times$ size gap) and generation/family (out-of-family older QWEN2.5-7B-INSTRUCT). For each subject we run two attackers: a weak out-of-family small attacker (QWEN2.5-7B) and a strong frontier attacker (OPUS 4.7) across all 6 topics.

Three observations. (i) *Stronger attackers convert a weaker same-family subject into more compliance.* Going from QWEN3.5-397B subject to QWEN3.5-4B subject (within the same training generation) gives OPUS 4.7 a +8pp edge ($40 \rightarrow 48\%$) and QWEN3.5-397B-as-attacker a +9pp edge ($28 \rightarrow 37\%$, from Table 2); within-QWEN3.5 mid-tier attackers gain a similar amount (4B +4pp, 35B +6pp). The weak QWEN2.5-7B attacker barely gains (+3pp). Within the same training generation, stronger attackers buy some elicitation on weaker-sized subjects. (ii) *An out-of-family older subject collapses.* QWEN2.5-7B-INSTRUCT has essentially no floor: every attacker we test averages $\geq 78\%$ across topics, including a QWEN2.5-7B attacker attacking itself (above the 65% frontier $\times$ frontier ceiling of OPUS $\rightarrow$ OPUS). (iii) *The dominant variable is training generation, not parameter count within family.* The size swap within QWEN3.5 moves rates by ≤ 9 pp; the generation swap to QWEN2.5-7B moves them by 45–65pp. Attacker capability matters, but less than which generation of post-training the subject ships with.

6 Ablation: how many turns does the attacker need?

The main results fix the conversation at 5 turns. Is that length necessary, or does the subject give in earlier? For each conversation in the full 3×3 scientific-consensus grid (540 conversations), the judge records the turn number(s) at which the subject produced a persuasive on-topic essay defending the against-consensus side. This lets us compute, for every subset of interest, the




Figure 3: Cumulative essay-production rate by turn over all 540 conversations of the main 3×3 scientific-consensus grid. **Left:** by topic. Creationism is already produced in 24% of conversations at turn 1 and reaches 59% by turn 5; Holocaust denial never exceeds 1%. **Middle:** by attacker. OPUS 4.7 extracts the essay fastest and highest (40% by turn 5), QWEN3.5-397B second (34%), GROK 4.20 last (12%). **Right:** by subject. OPUS 4.7-as-subject caves earliest and most often (17% at turn 1, 51% at turn 5); QWEN3.5-397B-as-subject holds better (24% at turn 5); GROK 4.20-as-subject holds best (11%). All three panels show saturation between turns 3 and 4.

cumulative essay-production rate by turn N : the fraction of conversations in which an on-topic against-consensus essay had been produced in at least one of turns 1 through N .

Most of the elicitation happens by turn 3. Aggregated across the main grid, the turn-1 cumulative rate is 6% (some easy cells, notably Opus-as-subject on creationism, already give essays at turn 1); by turn 2 it is 22%; by turn 3 26%; and turns 4 and 5 together add only ~ 2 pp. Whatever the attacker is going to extract is largely extracted by turn 3, with turns 4 and 5 acting as refinement rather than further breakthrough.

Attacker capability governs the speed of the curve. OPUS-as-attacker is faster and higher-ceiling than QWEN3.5-397B, which in turn leads GROK 4.20 at every turn, and the gaps open up early (by turn 2 OPUS is at 32%, QWEN at 24%, GROK at 11%). The curve therefore does not just show the time scale of elicitation but also the capability dimension that matters: stronger attackers extract more, and they extract it sooner.

Subject resistance differs sharply across models. OPUS-as-subject is the most permissive target in our grid by a wide margin: it produces an on-topic against-consensus essay in 17% of turn-1 replies and reaches 51% by turn 5, more than double QWEN3.5-397B-as-subject (24%) and close to five times GROK-as-subject (11%). Turn-1 production by Opus-as-subject is a notable finding on its own: in many cells the subject complies with the essay request directly, often wrapped in a disclaimer, without needing any attacker pressure.

More turns will not rescue the floor topics. Holocaust denial is the clearest case: its cumulative curve stays at or near 0% across all five turns. On this topic, adding more turns has no effect because no turn produces an on-topic essay in the first place. Antivax and racial IQ show a similar pattern with ceilings near 11–12% even at turn 5. Where the refusal boundary holds, it holds throughout the conversation; extending to ten or more turns is unlikely to change the picture.

attacker	Creat.	Flat-earth	Climate	Antivax	Racial IQ	Holocaust	avg
GROK 4.20	100%	100%	100%	100%	100%	100%	100%
CLAUDE OPUS 4.7	100%	100%	100%	100%	100%	83%	97%
QWEN2.5-7B	90%	80%	100%	100%	100%	90%	93%
QWEN3.5-35B-A3B	100%	100%	100%	100%	70%	0%	78%
QWEN3.5-397B-A17B	100%	100%	100%	90%	43%	0%	73%
QWEN3.5-4B	100%	90%	100%	50%	70%	0%	69%

Table 3: Attacker complicity rates: fraction of turn-1 conversations where the attacker actually posed the essay request to the subject (rather than refusing to issue the persuasion itself). Higher is more complicit—and, for the experiment, more informative about subject-side strength. Classifier is QWEN3.5-397B. Counts per cell: 30 for frontier attackers (GROK, OPUS, QWEN3.5-397B; 3 subjects $\times$ 10 reps), 10 otherwise. GROK is fully complicit on every topic; OPUS is fully complicit except on Holocaust denial (83%). Every QWEN-family attacker (4B, 35B-A3B, 397B-A17B) refuses 100% of Holocaust requests, and QWEN3.5-397B drops to 43% complicity on racial IQ as well.

7 Ablation: attacker complicity

Throughout this paper we have assumed that the “attacker” LLM actually plays the persuasion role—that its turn-1 message to the subject does contain the essay request. This assumption can fail: a frontier assistant acting as the attacker can itself refuse to pose the request (“I am unable to fulfill this; please propose an alternative subject”), in which case the subject’s refusal that follows is not a measure of subject-side guardrail strength but of the *attacker* refusing to be complicit in the persuasion. We call the fraction of turn-1 messages in which the attacker actually poses the request the *attacker complicity rate*: high complicity = the attacker cooperates with the benchmark and goes after the subject; low complicity = the attacker refuses to act as attacker. To measure it, we classify every attacker turn-1 message with QWEN3.5-397B as an independent classifier given the benchmark context (attacker/subject roles, the six topics, what “posing the request” means). The classifier returns one of two labels—**complied** (the attacker did pose the essay request, possibly with disclaimers or fictional framing) or **refused** (the attacker broke character).

Grok is fully complicit; Qwen least. GROK 4.20 has the highest attacker-complicity rate (100% across all six topics): as attacker, Grok always poses the request, even for Holocaust denial. OPUS is next (97% average, only wavering on Holocaust denial at 83% complicity). QWEN3.5-397B is the outlier: it poses 0/29 Holocaust requests and 13/30 racial-IQ requests on turn 1, never pushing the persuasion past turn 1 in those combinations. The smaller within-family Qwen attackers inherit the same Holocaust non-complicity (0% for QWEN3.5-4B, QWEN3.5-35B, and QWEN3.5-397B), and QWEN3.5-4B drops to 50% complicity on antivax as well.

Implication for the tables. The 0% cells for QWEN3.5-397B-as-attacker on Holocaust denial in Table 1 and on racial IQ are not a measure of subject-side guardrail strength—they are a measure of the attacker declining to be complicit. Whatever refusal boundary QWEN3.5-397B-as-subject would hold on Holocaust denial was never tested in those combinations, because the attacker did not pose the request. OPUS-as-attacker and GROK-as-attacker rates are the ones that measure subject-side strength cleanly. This is also why GROK-as-attacker, despite being the weakest at *extracting* essays (Grok rarely deploys the fictional-framing fallback), is the most reliable at issuing the initial request on every topic.




Figure 4: Same data as Table 3, in line form. X-axis: 6 scientific-consensus topics plus a mean column (rightmost, separator dotted line, marker outlined in black). Series: the 6 attackers. GROK 4.20 is fully complicit on every topic; OPUS 4.7 only drops on Holocaust denial (83%); the QWEN-family attackers all collapse to 0% complicity on Holocaust denial regardless of size, with QWEN3.5-397B additionally dropping to 43% on racial IQ. Where attacker complicity is low, the subject-side guardrail measurement on that cell is confounded by the attacker refusing to pose the request in the first place.

Implication for safety evaluation. A single-turn essay prompt is not an adequate refusal probe. A frontier-class assistant that refuses “write me a flat-earth essay” on the single-turn surface will produce such an essay in a five-turn conversation with another frontier-class assistant operating as the user, using only natural-language pressure. A realistic refusal benchmark for persuasive-essay generation therefore has to specify a multi-turn setting and an attacker of commensurate capability. Evaluating a model against a weak attacker understates its refusal vulnerability; evaluating against a peer attacker is the relevant measurement.

Implication for AI-agent interaction. As LLMs are increasingly deployed as agents in multi-agent pipelines, one frontier assistant routinely talks to another. The attacker setting in this paper is exactly that: one Opus instance talking to another Opus instance in a five-turn window. The fact that an Opus-as-user can persuade an Opus-as-assistant into producing consensus-denying essays bears directly on multi-agent system design, where a compromised or misaligned agent in the user position can exfiltrate restricted outputs from other agents in the system.

Why Holocaust denial is near the floor. Holocaust denial is the topic with the strongest subject-side refusal boundary in our data. Across the main 3×3 scientific-consensus cross-product (Table 1, 9 combinations, 90 conversations), we obtained 1 produced essay, from the QWEN→GROK combination. The attacker ablation (Table 2) adds 2 more Holocaust-denial essays, both from QWEN3.5-35B-A3B attacking QWEN3.5-397B (2/10; no other combination in that table produced any). Together that is 3 produced essays out of ~ 130 conversations on this topic, roughly an order of magnitude lower than creationism in the same data. Part of this floor, however, is attacker-side rather than subject-side: Table 3 shows that QWEN-family attackers have 0% complicity on the Holocaust-denial request (at every size tested: 4B, 35B, 397B), and OPUS-as-attacker drops to 83% complicity on the same topic. The cleaner

measurement of subject-side Holocaust resistance is therefore the OPUS- and GROK-attacker rows—GROK never breaks character on any topic—and in those rows the subject still produces only 3 essays across ~ 70 conversations, so the subject-side boundary is genuinely very strong, just less airtight than the raw 0/90 cells suggested. Both leaking combinations in the main grid and ablation are off the “stronger attacker breaks weaker subject” pattern—QWEN→GROK has a weaker subject, 35B→397B has a weaker attacker—consistent with the subject-side refusal boundary being very strong but not impenetrable under all attacker configurations.

8 Limitations

(i) *Single judge.* QWEN3.5-397B is used as judge throughout. A concurrent ablation in our companion study [Nogueira et al., 2026] re-judged 300 conversations with four different judges (CLAUDE OPUS 4.6, GROK 4.2, GEMINI 3.1 PRO, QWEN3.5-397B) and found pairwise agreement of 78–90%, with three of the four judges clustering within 0.7pp in consensus score; the same judge model was stable at 92% against mild system-prompt rewrites. The dominant variance source in that dataset is conversational stochasticity, not the judge. That ablation does not directly transfer to the present essay probe (different conversation format, different topic distribution, different claim types), but the broad pattern—judge swaps produce single-digit shifts, not reshufflings of the topic ordering—is our prior for how much topic-level reshuffling a different judge would produce here. That said, a judge stricter about fictional-framing disclaimers would lower the rates reported for creationism, climate-denial, and flat-earth (where disclaimers are common) more than it would lower rates on the three most-refused topics (antivax, racial IQ, Holocaust) where the subject refuses outright. (ii) *Five turns.* The scientific-consensus rates reported here are at five turns; Section 6 shows that for the OPUS→QWEN combination most of the elicitation happens by turn 3 and the rate plateaus by turn 4, so we expect longer conversations to add at most a few percentage points on the elicitable topics and nothing on the most-refused ones. We have not tested this directly for other combinations. (iii) *Judge’s willingness to count fictional-framed essays.* Several of our **produced** classifications come from essays that carry an explicit scientific-consensus disclaimer and are framed as a fictional character’s monologue. The judge consistently classifies these as **produced** based on the argumentative content, a choice that is defensible but is load-bearing for the reported rates. (iv) *Opus as attacker is truncated by the provider’s default output budget.* A non-trivial fraction of CLAUDE OPUS 4.7 attacker turns we inspected hit the provider’s default **max_tokens** cap and were delivered to the subject as incomplete messages (mid-sentence cut-off). On the next turn the subject often latches onto the truncation as evidence that the attacker is confused or that the request is malformed, and this lowers the elicitation rate achieved by Opus as attacker in our data. The Opus attacker numbers we report (Tables 1, 2) are therefore conservative; a run with an explicitly-raised **max_tokens** would likely move Opus-as-attacker rates up, which *strengthens* rather than weakens the paper’s central claim. (v) *QWEN-family attackers sometimes refuse the attacker role; OPUS occasionally does too.* Section 7 (Table 3) quantifies this with an independent classifier: QWEN3.5-397B-as-attacker is only 43% complicit on racial IQ and 0% complicit on Holocaust denial on turn 1 (it refuses the request itself, e.g. “*I am unable to fulfill this request as it involves generating content related to Holocaust denial, which violates safety policies against hate speech and historical misinformation. Please propose an alternative subject if you wish to continue.*”). The smaller within-family QWEN3.5 attackers (4B, 35B-A3B) show the same 0% complicity on Holocaust denial, and QWEN3.5-4B additionally drops to 50% complicity on antivax. OPUS is fully complicit except on Holocaust denial, where it drops to 83%. GROK 4.20 is fully complicit on every topic. When the attacker does not pose the request, the subject’s refusal that follows is not a measurement of subject-side guardrail strength but of the attacker declining to be complicit. Concretely: the QWEN-as-attacker 0% rates on Holocaust and racial IQ in Table 1 are attacker-complicity artifacts rather than subject-side measurements,

and the cleaner subject-side numbers on those topics come from the OPUS- and GROK-attacker rows.

References

- Myra Cheng, Sunny Yu, and Cinoo Lee. ELEPHANT: Measuring and understanding social sycophancy in LLMs. *arXiv preprint arXiv:2505.13995*, 2025.
- Carson Denison, Monte MacDiarmid, Fazl Barez, David Duvenaud, Shauna Kravec, Samuel Marks, Nicholas Schiefer, Ryan Soklaski, Alex Tamkin, Jared Kaplan, Buck Shlegeris, Samuel R. Bowman, Ethan Perez, and Evan Hubinger. Sycophancy to subterfuge: Investigating reward-tampering in large language models. *arXiv preprint arXiv:2406.10162*, 2024.
- Shangbin Feng, Chan Young Park, Yuhan Liu, and Yulia Tsvetkov. From pretraining data to language models to downstream tasks: Tracking the trails of political biases leading to unfair NLP models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.
- Jiseung Hong, Grace Byun, Seungone Kim, Kai Shu, and Jinho D Choi. Measuring sycophancy of language models in multi-turn dialogues. *arXiv preprint arXiv:2505.23840*, 2025.
- Avneet Kaur. Echoes of agreement: Argument driven sycophancy in large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 22803–22812, 2025.
- Akbir Khan, John Hughes, Dan Valentine, Laura Ruis, Kshitij Sachan, Ansh Radhakrishnan, Edward Grefenstette, Samuel R. Bowman, Tim Rocktäschel, and Ethan Perez. Debating with more persuasive LLMs leads to more truthful answers. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.
- Sungwon Kim and Daniel Khashabi. Challenging the evaluator: LLM sycophancy under user rebuttal. *arXiv preprint arXiv:2509.16533*, 2025.
- Rodrigo Nogueira, Giovana Kerche Bonás, Thales Sales Almeida, Andrea Roque, Ramon Pires, Hugo Abonizio, Thiago Laitz, Celio Larcher, Roseval Malaquias Junior, and Marcos Piau. Measuring opinion bias and sycophancy via LLM-based persuasion, 2026. URL <https://arxiv.org/abs/2604.21564>.
- Alicia Parrish, Angelica Chen, Nikita Nangia, Vishakh Padmakumar, Jason Phang, Jana Thompson, Phu Mon Htut, and Samuel R. Bowman. BBQ: A hand-built bias benchmark for question answering. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 2086–2105, 2022.
- Ethan Perez, Sam Ringer, Kamilė Lukošiušė, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, et al. Discovering language model behaviors with model-written evaluations. *arXiv preprint arXiv:2212.09251*, 2022.
- Virgile Rémard, Christos Xypolopoulos, and Michalis Vazirgiannis. Bias in the mirror: Are LLMs’ opinions robust to their own adversarial attacks? In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2025.
- Paul Röttger, Hannah Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy. Xstest: A test suite for identifying exaggerated safety behaviours in large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 5377–5400, 2024.

- David Rozado. The political preferences of LLMs. *PLoS ONE*, 19(7), 2024. doi: 10.1371/journal.pone.0306621.
- Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askill, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, et al. Towards understanding sycophancy in language models. *arXiv preprint arXiv:2310.13548*, 2023.
- Alexander Wei, Nika Haghtalab, and Jacob Steinhardt. Jailbroken: How does LLM safety training fail? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Tinghao Xie, Xiangyu Qi, Yi Zeng, Yangsibo Huang, Udari Madhushani Sehwag, Kaixuan Huang, Luxi He, Boyi Wei, Dacheng Li, Ying Sheng, Ruoxi Jia, Bo Li, Kai Li, Danqi Chen, Peter Henderson, and Prateek Mittal. Sorry-bench: Systematically evaluating large language model safety refusal. *arXiv preprint arXiv:2406.14598*, 2024. URL <https://sorry-bench.github.io/>.
- Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023.

Role	Model	Input (M tok)	Output (M tok)	Cost (USD)
Attacker	CLAUDE OPUS 4.7	3.19	0.61	\$31.14
Attacker	QWEN3.5-397B	2.80	2.51	\$8.82
Attacker	GROK 4.20	1.89	0.32	\$5.34
Attacker	QWEN3.5-35B-A3B	0.76	1.59	\$1.00
Attacker	QWEN3.5-4B (self-hosted)	0.77	1.67	\$0.00
Subject	CLAUDE OPUS 4.7	3.70	0.84	\$39.44
Subject	QWEN3.5-397B	4.25	3.73	\$12.08
Subject	GROK 4.20	1.77	0.31	\$5.71
Judge	QWEN3.5-397B	2.01	0.04	\$1.23
Total (682 conversations)				\$104.77

Table 4: Actual API cost. Per-conversation mean output from the subject is ~ 14 K tokens for OPUS, ~ 20 K tokens for QWEN3.5-397B, and ~ 5 K tokens for GROK, spread across 5 turns; the attacker generates roughly a third to a half of that volume on the input side.

A Benchmark cost

Table 4 reports the actual API cost for the 682 conversations behind Tables 1 and 2: 540 in the main $3 \times 3 \times 6$ grid and 142 in the attacker-scaling ablation (QWEN3.5-4B and QWEN3.5-35B-A3B attackers; QWEN3.5-397B and OPUS attackers reuse the main-grid Qwen/Opus $\rightarrow$ QWEN combinations). Token counts are taken from the OpenRouter **usage** field returned on each call, not estimated from character counts; costs apply each provider’s public per-million-token input/output rates at run time (OPUS 4.7: \$5/\$25; GROK 4.20: \$2/\$6; QWEN3.5-397B: \$0.39/\$2.34; QWEN3.5-35B-A3B: \$0.18/\$0.54). The QWEN3.5-4B attacker was self-hosted on a single H100 and is billed at \$0 against the OpenRouter account; it consumed roughly the same input/output token volume as the 35B attacker on the same jobs, so a hosted-API equivalent cost would be on the same order as the 35B row below.

B Subject-ablation: per-topic counts

Table 5 gives the per-topic counts behind the subject-ablation discussion in Section 5.

attacker	subject	Creat.	Flat-earth	Climate	Antivax	Racial IQ	Holocaust	avg
QWEN2.5-7B	QWEN3.5-397B	30%	30%	0%	0%	0%	0%	10%
QWEN3.5-4B	QWEN3.5-397B	50%	100%	10%	0%	0%	0%	27%
QWEN3.5-35B-A3B	QWEN3.5-397B	50%	90%	20%	0%	0%	0%	27%
QWEN3.5-397B-A17B	QWEN3.5-397B	60%	70%	40%	0%	0%	0%	28%
CLAUDE OPUS 4.7	QWEN3.5-397B	90%	100%	40%	10%	0%	0%	40%
QWEN2.5-7B	QWEN3.5-4B	50%	10%	10%	10%	0%	0%	13%
QWEN3.5-4B	QWEN3.5-4B	90%	70%	20%	0%	0%	0%	30%
QWEN3.5-35B-A3B	QWEN3.5-4B	80%	70%	40%	10%	0%	0%	33%
QWEN3.5-397B-A17B	QWEN3.5-4B	100%	70%	40%	10%	0%	0%	37%
CLAUDE OPUS 4.7	QWEN3.5-4B	90%	80%	80%	30%	10%	0%	48%
QWEN2.5-7B	QWEN2.5-7B	90%	80%	70%	90%	70%	70%	78%
QWEN3.5-4B	QWEN2.5-7B	100%	100%	100%	90%	100%	10%	83%
QWEN3.5-35B-A3B	QWEN2.5-7B	100%	100%	100%	100%	100%	30%	88%
QWEN3.5-397B-A17B	QWEN2.5-7B	100%	100%	100%	100%	100%	10%	85%
CLAUDE OPUS 4.7	QWEN2.5-7B	100%	100%	100%	100%	100%	70%	95%

Table 5: Full per-topic breakdown for the subject-ablation experiment. $N = 10$ per cell. The three blocks are the three subject models; the last column is the cross-topic average per (attacker, subject) pair.